

Myar Archiving Program (due 9/27/2020)

Learning Objectives

Upon completion of this assignment, you should be able:

1. Be able to efficiently use low level fileio system calls
2. Understand and use the metadata that describes a file
3. Be able to interact with directory structures
4. Interact with system data structures compatibly (ar_hdr, archive format)

New mechanisms you will see and use include:

- system calls: read, write, lseek, stat, utime, unlink
- library routines: Directory Routines, String Routines (malloc, free only for buffer size from stat)
- umask

NAME

`myar` - Archiving Utility

SYNOPSIS

`myar` [qxotvd:A:] archive-file [file1]

DESCRIPTION

`myar` operates on the archive-file listed adding, deleting, listing contents according to the operations selected, using the files specified as appropriate. `Myar` maintains compatibility with “`ar`” utility on Linux, with the exception of the “`A`” option which is an added option to `myar`, and any item mentioned in the notes below.

- q Quickly append the files specified to the archive-file
- x Extract the specified files from the archive-file
- o Used in combination with “`x`” restore the original permission and mtime
- t List the file names in the archive-file
- v Used in combination with “`t`” print the verbose info as “`ar`” does.
- d Delete the files specified from the archive-file
- A Quickly append all regular files in the directory modified more than N days ago

Exit Status


`myar` returns 0 on success and 1 on failure

ERROR MESSAGES

Same as in `ar`.

DISCUSSION

This is what an AR Archive Files looks like



The archive file maintained must use exactly the standard format defined in `"/usr/include/ar.h"`, and in fact must be tested with archives created with the `"ar"` command and `ar` must be tested with archives created by `myar`. You must

```
#include <ar.h>
```

for the definition of `struct ar_hdr`. In addition, define the following

```
struct meta {  
    char name[16]; //room for null  
    int mode;  
    int size;  
    time_t mtime; // a time_t is a long  
}
```

and write functions

```
int fill_ar_hdr(char *filename, struct ar_hdr *hdr);
```

```
int fill_meta( struct ar_hdr hdr, struct meta *meta);
```

The first of these `"stat"`s the filename, and fills in the `ar_hdr`. `stat` structures have binary fields, `ar_hdr` fields are all printable ascii, so conversions are necessary. This is used in adding members to the archive.

The second goes in the opposite direction, converting fields in the `ar_hdr` into binary for the meta structure. The meta structure is used when extracting members from the archive, and for the `"-t -v"` option.

NOTES

- This is an exercise in efficiently using `read/write/lseek` directly, and you cannot use the buffered I/O library (`fopen`, `fread`, `putc`, etc). You should be reading and writing `ar_hdrs` in a single system call. File contents is written in buffers sized by `stat`.
- You may assume all files are in the current directory, no filename is longer than 15 bytes, and no archive index will be attempted. Unlike `"ar"` you will not follow symbolic links. In the `"A"` option you skip symbolic links.
-
- The `ar_hdr` and the `ARMAG` have even size. `ar_hdr`'s must always be on an even boundary. This is maintained if the file being archived is even. However when archiving a file whose size is odd, a padding byte is required after it to preserve `ar_hdr` alignment.
- You MUST NOT consult source code for other implementations of `"ar"`
- Create archives with permission 0666 subject to `umask`. Extraction permission in the `-x -o` option is also subject to the `umask`. (Good news, this means there is nothing to do).
- When running the `"ar"` command on Ubuntu use `"qU"` instead of `"q"`. Ubuntu does not include the metadata without the `"U"`. This is a recent optimization justified by the typical use of `"ar"` to maintain binary libraries. `Myar` should always include the metadata.
- The `"x"` and `"d"` commands operate on the first file matched in the archive, without checking for further matches.
- In the case of the `"d"` option, you will have to build a new archive file to recover the space. Do this by unlinking the original file after it is opened, and creating a new archive with the original name.
- You are required to handle as many files as appear on the command line. Since file I/O is expensive, do not make more than one pass through the archive file, an issue especially relevant to the multiple delete case.

HAND in myar.c, Makefile, any other files (not ar.h), README in a lab2 directory created in <yourid>/cs551